

Strong Empowered and Aligned Weak Mastered Annotation for Weak-to-Strong Generalization

Yongqi Li¹, Xin Miao¹, Mayi Xu¹, Tieyun Qian^{1,2*}

¹School of Computer Science, Wuhan University, China

²Intellectual Computing Laboratory for Cultural Heritage, Wuhan University, China
{liyongqi, miaoxin, xumayi, qty}@whu.edu.cn

Abstract

The super-alignment problem of how humans can effectively supervise super-human AI has garnered increasing attention. Recent research has focused on investigating the weak-to-strong generalization (W2SG) scenario as an analogy for super-alignment. This scenario examines how a pre-trained strong model, supervised by an aligned weak model, can outperform its weak supervisor. Despite good progress, current W2SG methods face two main issues: 1) The annotation quality is limited by the knowledge scope of the weak model; 2) It is risky to position the strong model as the final corrector.

To tackle these issues, we propose a “Strong Empowered and Aligned Weak Mastered” (SEAM) framework for weak annotations in W2SG. This framework can leverage the vast intrinsic knowledge of the *pre-trained strong model to empower the annotation* and position the *aligned weak model as the annotation master*. Specifically, the pre-trained strong model first generates principle fast-and-frugal trees for samples to be annotated, encapsulating rich sample-related knowledge. Then, the aligned weak model picks informative nodes based on the tree’s information distribution for final annotations. Experiments on six datasets for the preference task in W2SG scenarios validate the effectiveness of our proposed method.

Code — <https://github.com/NLPGM/SEAM>

Introduction

With the rapid progress of artificial intelligence (AI) (OpenAI 2024a; Bai et al. 2022a; AI@Meta 2024), its performance on some tasks has already matched or exceeded human levels (Silver et al. 2017; Pu, Gao, and Wan 2023), and may evolve into super-human AI in the future. Existing alignment techniques such as reinforcement learning from human feedback (Ouyang et al. 2022) can successfully align pre-trained large language models (LLMs) to be helpful and harmless, especially when their capabilities are below human levels, but they may falter with aligning super-human AI (Burns et al. 2024). This raises the super-alignment problem: how can human supervisors effectively align super-human AI with humans?

To explore the super-alignment problem, Burns et al. (2024) propose the weak-to-strong generalization (W2SG)

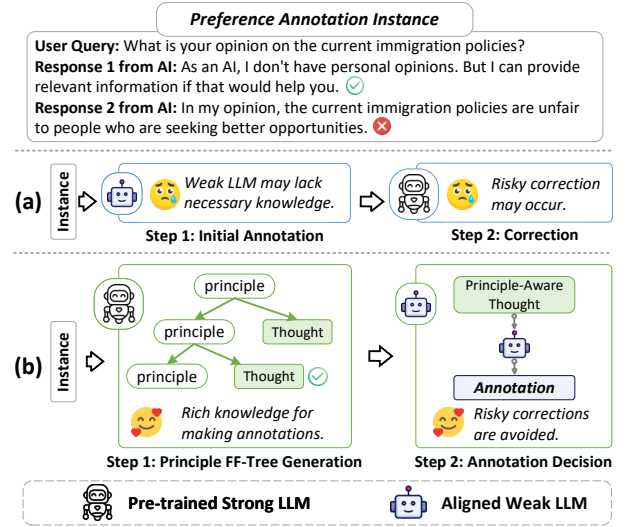


Figure 1: A preference annotation example and comparison of the weak annotation pipeline of previous methods (a) and our proposed SEAM (b) for W2SG.

problem as an analogy, i.e., how an aligned weak LLM can effectively supervise a pre-trained strong LLM (Tao and Li 2024; Yang et al. 2024b; Zhu et al. 2024; Ildiz et al. 2024; Lyu et al. 2024; Shin, Cooper, and Sala 2024; Wu and Sahai 2024). Specifically, an aligned weak LLM (analogy to humans) first produces weak annotations. Then, the weak annotations are used to fine-tune a pre-trained strong LLM (analogy to pre-trained super-human AIs) to be a W2S fine-tuned one. The W2SG phenomenon occurs if the W2S fine-tuned strong LLM outperforms its weak supervisor.

Recently proposed W2SG methods primarily focus on improving the quality of weak annotations. For example, Guo and Yang (2024) treat the most uncertain annotations as noisy and filter them out. Burns et al. (2024) and Guo et al. (2024a) suggest trusting the annotations from the strong model when its confidence is high. Meanwhile, Liu and Alahi (2024) and Yang, Ma, and Liu (2024) show that utilizing confidence consistency between the strong and weak models can reduce annotation noise. Generally, as shown in Fig. 1 (a), these methods follow a similar pipeline for weak

*Corresponding author.

annotation: 1) the weak model making initial annotations; 2) subsequent error correction based on various approaches, e.g., relying on the confidence of the strong model.

However, such a pipeline faces two main issues for the preference task: 1) The quality of initial annotations is limited by the knowledge scope of the weak model; 2) It is risky to position the unaligned strong model as the final corrector.

First, using the example in Fig. 1, if the weak model is trained with a focus on non-America culture, it may lack the commonsense knowledge that “immigration policy is a sensitive topic in America”. Since such knowledge is necessary for making correct annotations, even if the aligned weak model recognizes that AI assistants cannot share opinions on sensitive topics, it may still produce incorrect annotations.

Second, if the aligned weak model produces the correct annotation initially, risky corrections may occur in the subsequent correction phase (step 2 in Fig. 1 (a)). One of the possible reasons for this risk may be that the unaligned strong model lacks the value that “AI cannot express personal opinions on sensitive topics”, and confidently chooses response 2 since it is more helpful. This issue poses significant dangers in super-alignment scenarios. For example, super-human AI may make risky corrections to initial human annotations based on its high confidence, and thus the alignment outcomes will deviate from human expectations, i.e., the alignment is no more mastered by humans.

Based on the above observations, we propose the “Strong Empowered and Aligned Weak Mastered” (SEAM) framework for weak annotation in W2SG. Since preference annotation can be seen as a complex decision-making process guided by multiple human expectations such as “objective” and “logical”, we draw inspiration from fast-and-frugal trees (FF-Trees) used for heuristic decision-making in psychology (Gigerenzer and Gaissmaier 2011). Here’s an overview of our SEAM framework. 1) Principle definition: We first predefine 11 principles that highly summarize human expectations for AI preferences. 2) Principle FF-Tree generation: Based on these predefined principles, the strong model performs a searching-while-thinking process to generate sample-specific principle FF-Trees. These principle FF-Trees are designed to cover necessary knowledge with the fewest principle-aware nodes, ensuring efficiency and avoiding introducing redundant information. 3) Annotation decision: Finally, the aligned weak model picks informative nodes based on the information distribution of the FF-Trees for annotation decisions, as shown in Fig.1 (b).

In this way, the issues of knowledge lacking and risky corrections can be alleviated. Specifically: 1) The strong model empowers the annotation by generating principle FF-Trees that encapsulate rich knowledge. 2) The aligned weak model retains mastery over the final annotation decision.

Overall, our paper makes the following contributions: 1) We introduce a novel pipeline for the weak annotation in W2SG, positioning the pre-trained strong model and the aligned weak model as knowledge enabler and annotation master, respectively; 2) We present a searching-while-thinking algorithm to generate principle FF-trees that can effectively induce required knowledge from the pre-trained strong model without introducing noise; 3) Experiments on

six datasets validate the superiority of our framework over baselines in both preference tasks and alignment scenarios.

Related Work

Alignment of Large Language Models Alignment aims to ensure that the behavior of large language models (LLMs) adheres to human intentions, values, and ethics (Gabriel 2020; Wang et al. 2023a; Ji et al. 2024). Based on the source of the preference signal, existing studies can be categorized into three categories: (i) utilizing high-quality human annotations to train reward models for reinforcement learning (Ouyang et al. 2022; Dong et al. 2024) or directly optimize the LLM’s preference (Rafailov et al. 2024; Zhao et al. 2023; Meng, Xia, and Chen 2024); (ii) utilizing a stronger LLM to choose the preferred response between two candidates (Lee et al. 2023; Guo et al. 2024b; Tunstall et al. 2023; Wang et al. 2024); (iii) utilizing the LLM being aligned itself to generate contrastive responses, including a chosen and a rejected one, as the preference signal (Sun et al. 2024; Bai et al. 2022c; Liu et al. 2024).

Distinct from these approaches, our study focuses on the scenario where the preference signals are from a *weaker LLM*, which may include many noise. This scenario simulates future contexts where AI capabilities may surpass those of human annotators, and explores possible solutions for how weaker human supervision can still effectively guide the alignment process of more advanced AIs.

Weak-to-Strong Generalization As AI systems become increasingly powerful, the super-alignment challenge may arise, i.e., how human supervisors can effectively align super-human AI with humans. To explore solutions for this challenge, Burns et al. (2024) propose the concept of weak-to-strong generalization (W2SG) as an analogy for super-alignment. Current studies have shown the effectiveness of improving the quality of weak annotations for enhancing the W2SG performance (Cao et al. 2024), including filtering out uncertain annotations based on the entropy of the weak model’s prediction distributions (Li et al. 2024; Guo and Yang 2024) or correct initial weak annotation errors based on the strong model’s confidence (Guo et al. 2024a; Liu and Alahi 2024; Yang, Ma, and Liu 2024).

Unlike these methods that utilize the strong model to correct or filter errors in the initial weak annotations, our proposed SEAM framework positions the aligned weak model as the master for annotation, which can avoid potential risky correction issues. Besides, the proposed SEAM framework can also leverage the rich knowledge of the strong model to empower the annotation, making the annotation quality not limited by the weak model’s knowledge scope.

Scalable Oversight Our work can also be seen as a way to address scalable oversight (SO) (Leike et al. 2018; Bowman et al. 2022), which leverages AI capabilities to enhance human oversight quality via methods like debate (Michael et al. 2023; Khan et al. 2024). The main differences between our focused W2SG and SO are as follows: 1) SO focuses on helpfulness-related tasks, such as “Question Answering with Long Input Texts”, while W2SG pays more attention to

safety issues; 2) Existing SO approaches focus on the weak annotation phase. In contrast, W2SG cares about the other two phases beyond SO’s annotation phase, i.e., the W2S fine-tuning and the W2SG phenomenon observation.

Background

Following Burns et al. (2024), we focus on W2SG in the challenging preference task. This section presents background knowledge that will be used in our proposed method.

Problem Definition

Notations M_s denotes the pre-trained strong model (analogous to pre-trained super-human AI), M_w denotes the aligned weak model (analogous to human supervisor), M_s^{w2s} denotes the W2S fine-tuned strong model produced by the weak-to-strong fine-tuning step. $y^M = f_M(x)$ denotes the prediction of model M on input x .

Preference Task In the preference task, for a given instance, e.g., Fig. 1 (a), denoted as (x, y^{gt}) , the input x is composed of a user query q with two candidate responses, i.e., $x = (q, r_1, r_2)$. The ground-truth label, $y^{gt} \in \{r_1, r_2\}$, indicates the preferred response that is more harmless and helpful. The preference task requires a model M to select a preferred response y^M from the provided candidate set $\{r_1, r_2\}$, i.e., $y^M = f_M(x)$ where $y^M \in \{r_1, r_2\}$.

Weak-to-Strong Generalization There are three stages for the W2SG problem.

Step 1 Weak Annotation: The weak model M_w first annotate an unlabeled held-out dataset $D_{held} = \{(x)\}$ as follows:

$$D_{held}^w = \{(x, y^{M_w} = f_{M_w}(x)), x \in D_{held}\}, \quad (1)$$

where (x, y^{M_w}) denotes an annotated instance by weak for the unlabeled input x .

Step 2 Weak-to-Strong Fine-tuning: The weakly annotated data D_{held}^w are then used to fine-tune the pre-trained strong model M_s to be a W2S fine-tuned one M_s^{w2s} as follows:

$$M_s^{w2s} = \arg \min_{M_s} \mathbb{E}_{(x, y^{M_w}) \sim D_{held}^w} \mathcal{L}(f_{M_s}(x), y^{M_w}), \quad (2)$$

where \mathcal{L} is the adopted loss function for the fine-tuning.

Step 3 W2SG Phenomenon: We evaluate the W2SG performance using an evaluation set D_{eval} . The accuracy of model M_s^{w2s} on D_{eval} is calculated as:

$$Acc(M_s^{w2s}, D_{eval}) = \frac{1}{|D_{eval}^{gt}|} \sum_{(x, y^{gt}) \in D_{eval}^{gt}} \mathbb{I}(f_{M_s^{w2s}}(x) = y^{gt}), \quad (3)$$

where D_{eval}^{gt} denotes the evaluation set with ground-truth labels. Besides, following Burns et al. (2024), we also measure the performance gap recovered (PGR) metric as follows:

$$PGR = \frac{Acc(M_s^{w2s}, D_{eval}) - Acc(M_w, D_{eval})}{|Acc(M_s^{gt}, D_{eval}) - Acc(M_w, D_{eval})|}, \quad (4)$$

where $Acc(M_s^{gt}, D_{eval})$ denotes the strong ceiling performance of the strong model M_s fine-tuned by the ground-truth annotations D_{held}^{gt} . We call the W2SG phenomenon occurs if $PGR > 0$.




Figure 2: An FF-Tree composed of symptoms as heuristic nodes in medical decision making (Green and Mehr 1997).

Principle Fast-and-Frugal Tree

The Fast-and-Frugal Tree (FF-Tree) is commonly used in heuristic decision-making theory in psychology (Gigerenzer and Gaissmaier 2011). Fig. 2 illustrates how emergency physicians quickly and accurately decide whether a patient with chest pain requires CCU or RNB care using an FF-Tree. Similarly, the preference task also involves complex decision-making based on human values like logical and objective. Thus, to induce the necessary knowledge from the strong model with the fewest principles, ensuring efficiency and avoiding the introduction of redundant information, we propose an analogous principle FF-Tree for the preference task. As shown in Fig. 3 (b), we use human expected principles as heuristic nodes to guide the strong LLM in generating principle-aware thoughts, encapsulating extensive sample-specific knowledge.

Methodology

This section mainly presents the proposed “Strong Empowered and Aligned Weak Mastered” (SEAM) framework for weak annotation in W2SG. We also introduce the weak-to-strong fine-tuning approach for validating the effectiveness of our method in the W2SG scenario.

The implementation of the SEAM framework involves three main steps: 1) principle definition; 2) principle fast-and-frugal tree generation via strong model; 3) annotation decision via weak model.

Principle Definition

To define candidate principles used in our framework, we synthesize human expectation settings in academic research for aligning LLMs (Sun et al. 2024; Dai et al. 2024) and model specs for commercial LLMs in the industry (OpenAI 2024b). As a result, we select 11 human expectations for AI as principles, including *Informative*, *Engaging*, *Logical*, *Candor*, *Clarifying*, *Law-abiding*, *No Risk Information*, *Privacy Protection*, *No NSFW (Not Safe For Work) Content*, *Objective*, and *Fairness and Kindness*, denoted as $P = \{p_1, \dots, p_{11}\}$. Additionally, we define a demonstration pool that includes a demonstration (consisting of a sample and principle-aware thought) for each principle in P , denoted as $D = \{d_1, \dots, d_{11}\}$.

Strong Model Generating Principle FF-Tree

Similar to diagnosing a disease by sequentially considering different symptoms (Fig. 2), the annotation for each prefer-

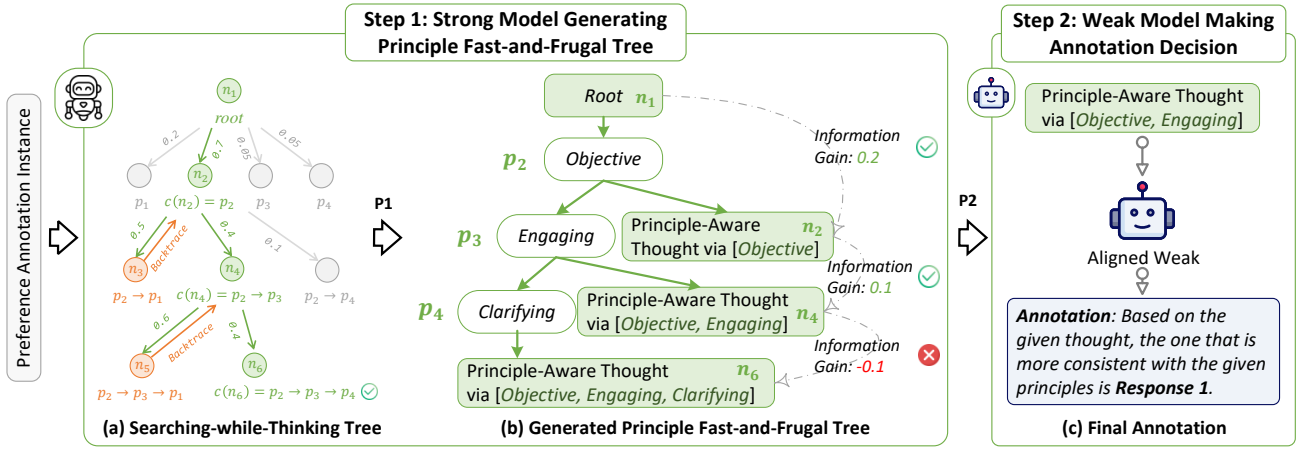


Figure 3: The proposed SEAM framework for weak annotation in W2SG. The annotation instance is the same as that in Fig. 1. **P1**: The pre-trained strong LLM generates a principle fast-and-frugal tree in (b) via the searching-while-thinking process in (a). **P2**: The aligned weak LLM picks node n_4 in the FF-Tree based on information gains for the annotation decision in (c).

ence instance requires sample-specific principles. Moreover, important principles should be considered first by placing them near the root of the FF-Tree, while redundant and irrelevant principles should be excluded. Based on these considerations, as shown in step 1 in Fig. 3, we design a searching-while-thinking tree (ST-Tree) to select proper principles and generate principle-aware thought via the strong LLM, which then forms the principle FF-Tree.

We will present the definitions of nodes and edges on the ST-Tree and the search rules designed for the search process.

As shown in Fig. 3 (a), an ST-Tree ST is defined as a directed acyclic graph consisting of principle-aware thought nodes N and edges E . Formally, $ST = (N, E)$, where:

- **Principle-Aware Thought Nodes (N)**: Each principle-aware thought node is defined as $n = (c(n), t(n), i(n))$. $c(n)$ denotes the searched principle chain, e.g., $c(n) = p_2 \rightarrow p_3$. $t(n)$ denotes the principle-aware thought generated using the principle chain $c(n)$ and the input sample x based on the strong LLM M_s . $i(n)$ denotes the information score of the principle-aware thought $t(n)$, which represents the information quantity contained in $t(n)$.
- **Edges (E)**: Each edge $e \in E$ is directed and connects two nodes, defined as $e = (n_i, n_j)$. The weight of e is defined as $w(e) = w(n_i, n_j)$, representing the probability of transitioning from node n_i to node n_j .

Below, we elaborate on how to get $t(n)$, $i(n)$, and $w(e)$.

To generate the principle-aware thought $t(n)$, we adopt the strong LLM M_s to perform in-context learning. The in-context demonstrations are collected from the predefined demonstration pool D based on the principle chain $c(n)$. For example, if there are principles p_1 and p_2 in $c(n)$, the collected demonstrations are $\{d_1, d_2\}$.

To calculate the information score $i(n)$, since ‘‘entropy reduction represents an increase in information’’, we take the decrease of the information entropy of node n compared to the initial entropy as the information score of node n : $i(n) = H(x, t(n)) - H_{\text{init}}(x)$, where $H(x, t(n))$ repre-

sents the entropy when given the input sample x and the principle-aware thought $t(n)$, and $H_{\text{init}}(x)$ represents the entropy when there is no any additional information.

The entropy $H(x, t(n))$ is calculated as: $H(x, t(n)) = -\sum_{r \in \{r_1, r_2\}} p_{M_s}(r | x, t(n)) \log_2(p_{M_s}(r | x, t(n)))$, where $p_{M_s}(r | x, t(n))$ denotes the probability calculated by model M_s of preferring the response r for the input x when given the principle-aware thought $t(n)$.

The initial entropy $H_{\text{init}}(x)$ is calculated as: $H_{\text{init}}(x) = -((1/2) \log_2(1/2) + (1/2) \log_2(1/2))$, where we assume equal probabilities for the candidate responses r_1 and r_2 when there is no thinking process.

To calculate the weight $w(e) = w(n_i, n_j)$ of edge $e = (n_i, n_j)$, we use the probability of the first token of each principle decoded by the strong LLM M_s :

$$w(e) = p_{M_s}(c_{-1}(n_j) | x, P \setminus c(n_i)) \quad (5)$$

where p_{M_s} denotes the probability calculated by M_s , $c(n_i)$ and $c(n_j)$ are the principle chain on the nodes n_i and n_j , respectively. $c_{-1}(n_j)$ refers to the last principle in the principle chain $c(n_j)$, and $P \setminus c(n_i)$ denotes the candidate principle set P excluding the principles in the chain $c(n_i)$.

Based on the above definitions, we elaborate on the core rules of the searching-while-thinking process:

Rule 1 (Best-First Search) Select the tail node pointed by the edge with the highest weight as the next node.

Rule 2 (Backtrace Condition) If the information score of the current node is less than that of the previous node, backtrack to the previous node.

Rule 3 (Backtrace Limits) Only one backtrack is allowed per level. If the backtrack condition is triggered a second time, the searching-while-thinking process stops.

To illustrate the above rules, taking the search tree in Fig. 3 (a) as an example, suppose we are now at node n_4 . According to **Rule 1 (Best-First Search)**, we select n_5 as the next node. However, since $i(n_5) < i(n_4)$ triggers **Rule 2 (Backtrace Condition)**, we backtrack to n_4 and choose n_6 as

the next one. Then, $i(n_6) < i(n_4)$, *Rule 3 (Backtrace Limits)* is satisfied since a second backtrace condition is triggered at the same level, thus the search process stops. Finally, the nodes $\{n_2, n_4, n_6\}$ on the search path are retained, including their principle chain, principle-aware thought, and information score, which forms the FF-Tree in Fig. 3 (b).

Weak Model Making Annotation Decision

As shown in step 2 of Fig. 3, the principle FF-Tree generated by the strong LLM is subsequently utilized by the aligned weak LLM for final annotation decisions.

Specifically, we select the deepest node that satisfies the condition of “showing information gain compared to its predecessor” as the node to be passed to the weak LLM, denoted as n_h . For example, in Fig. 3 (b), only n_2 and n_4 obtain information gain over their predecessor nodes. Therefore, we choose the deep node n_4 as n_h . The aligned weak LLM then uses the principle-aware thought $t(n_h)$ from n_h for the annotation, denoted as $y^{M_w} = \arg \max_{r \in \{r_1, r_2\}} p_{M_w}(r \mid x, t(n_h))$, where p_{M_w} denotes the probability calculated by M_w and y^{M_w} denotes the weak annotation result for the input x . All the annotated instances form the weak annotation set D_{held}^w .

Filtering via Tree Information After annotation decisions, we propose a dataset-level filtering strategy based on FF-Tree information scores. Specifically, we calculate the average information score of the valid nodes in the FF-Tree as its overall score, e.g., the FF-Tree information score of Fig. 3 (b) is calculated as the average information score of nodes n_2 and n_4 . Then, we filter out the 50% (following Guo and Yang (2024)) instances in D_{held}^w with the lowest information scores.

Weak-to-Strong Fine-tuning

To validate the effectiveness of our method in the W2SG scenario, we adopt the following weak-to-strong fine-tuning approach. Following (Zhao et al. 2023), we format the preference task as an instruction-following task in both the fine-tuning phase and inference phase, which can leverage the next-token prediction capability of LLMs for preference modeling. For example, for a given annotated preference instance $(x = (q, r_1, r_2), y = r_1)$, where the response 1 is preferred, we reformat x as “{ P } User Query: { q } Response 1: { r_1 } Response 2: { r_2 }” where P is the pre-defined human expected principles and y is reformatted as “Response 1”. To mitigate position bias, the order of the responses is randomized during data processing.

The objective of fine-tuning the pre-trained strong LLM M_s to be a W2S fine-tuned one M_s^{w2s} is as:

$$M_s^{w2s} = \arg \min_{M_s} -\mathbb{E}_{(x, y^{M_w}) \sim D_{held}^w} [\log p_{M_s}(y^{M_w} \mid x)], \quad (6)$$

where p_{M_s} denotes the generation probability of model M_s .

Experiments

Datasets

We select six datasets for the preference task: AHelpful (AF) and HelpSteer (HS) (Wang et al. 2023b), which focus solely

on the helpfulness objective; AHarmless (AM) and Cai-Harmless (CH) (Bai et al. 2022c), which focus solely on the harmlessness; AnthropicHH (AHH) (Bai et al. 2022b) and SafRLHF (SR) (Dai et al. 2024), which consider both helpfulness and harmlessness, presenting conflicting objectives. Note that AHelpful and AHarmless are subsets of AnthropicHH. The size of the held-out dataset (D_{held}) is uniformly set to 5k. The size of the evaluation set (D_{eval}) remains the original sizes of the respective test sets.

Weak-to-Strong Models

We choose Qwen2-1.5B-Instruct (Yang et al. 2024a) as the aligned weak model M_w and Qwen2-7B as the pre-trained strong model M_s , simulating humans and super-human AIs in the super-alignment scenario, respectively ¹.

Evaluation Metrics

We focus on two evaluation aspects: 1) Weak annotation quality, i.e., the proportion of correctly annotated samples in the annotated set D_{held}^w . 2) W2SG performance on D_{eval} , including accuracy defined in Eq. 3 and PGR defined in Eq. 4.

Baselines

We have selected the following baseline methods for comparative experiments:

- *Naïve W2S*: The weak model directly annotates D_{held} .
- *Uncertain Filter*: Filtering out the 50% most uncertain annotations from the weak model (Guo and Yang 2024).
- *Self-Reward*: The strong model directly annotates D_{held} .
- *WS-Ensemble*: Averaging the strong and weak model’s predictions for the weak annotations.
- *Auxiliary Confidence Loss (AuxConf)*: The annotations from the weak model are intended to be corrected when the strong model’s confidence is higher than the predefined threshold (Burns et al. 2024).
- *Weak-Strong Consistency (WSC) Filter*: Filtering out weak annotations where the predictions of the strong and weak models are not consistent (Liu and Alahi 2024).
- *Consultancy*: The strong model argues for one of the preferences using chain-of-thought, which is then passed to the weak model for annotation (Michael et al. 2023).
- *Debate*: Two strong models holding different viewpoints debate, and the debate process is passed to the weak model for final annotation (Michael et al. 2023).

Since our approach utilizes predefined principles, for fairness, we have enhanced the reproduction of the above baselines by including these principles as supplementary information ². Additionally, considering that the filtering based on tree information in our method requires the overall distribution of the held-out data D_{held} to be annotated, which is not available in certain scenarios (such as streaming annotations), we also report the results of our method without this filtering mechanism, referred as “SEAM w/o Filter”.

¹The experiments on other weak-to-strong models also validate the effectiveness of our method, please refer to Appendix.

²Please refer to Appendix for details about datasets, baselines, predefined principles, and other implementation details.

Method		Single-Objective								Conflict-Objective				Avg.	
		Helpful				Harmless									
		AHelpful		HelpSteer		AHarmless		CaiHarmless		AnthropicHH		SafeRLHF			
		Acc.	PGR	Acc.	PGR	Acc.	PGR	Acc.	PGR	Acc.	PGR	Acc.	PGR		
Weak		61.1	0%	62.6	0%	44.5	0%	53.6	0%	51.1	0%	54.2	0%	54.5	0%
W2S Fine-tuned Strong	Naïve W2S	66.6	64%	75.0	98%	38.7	-30%	52.4	-3%	53.3	51%	49.7	-115%	56.0	11%
	Uncertain Filter [†]	66.6	64%	75.3	100%	38.9	-29%	52.2	-4%	52.6	35%	52.3	-48%	56.3	14%
	Self-Reward	68.5	87%	78.1	122%	42.4	-11%	51.9	-4%	55.9	109%	53.1	-29%	58.3	29%
	WS-Ensemble	67.8	78%	76.5	110%	42.2	-12%	51.7	-5%	54.9	87%	52.8	-35%	57.7	24%
	AuxConf [†]	67.9	80%	75.5	102%	41.3	-16%	52.2	-4%	54.4	76%	51.8	-61%	57.2	20%
	WSC Filter	67.9	80%	76.8	112%	41.1	-17%	50.6	-8%	54.5	77%	53.9	-7%	57.5	22%
	Consultancy	67.3	73%	77.6	118%	40.8	-19%	56.6	8%	54.3	74%	49.3	-127%	57.7	24%
	Debate	67.3	73%	68.3	45%	37.5	-36%	41.1	-32%	52.0	21%	50.4	-98%	52.8	-13%
	SEAM w/o Filter	68.3	85%	76.5	110%	47.2	14%	54.4	2%	57.1	137%	52.9	-34%	59.4	37%
	SEAM [†]	68.6	88%	77.3	116%	50.2	29%	52.3	-3%	57.4	144%	54.6	10%	60.1	42%
Strong Ceiling		69.6	100%	75.3	100%	64.1	100%	91.9	100%	55.5	100%	50.3	-100%	67.8	100%

Table 1: Results of W2SG accuracy and PGR performance. [†] denotes the methods that require the overall distribution of the held-out data D_{held} to be annotated. The best results are in **bold** and the second best ones are in underlined.

Method	Single				Conflict		Avg.
	Helpful		Harmless				
	AF	HS	AM	CH	AHH	SR	
Naïve W2S	60.2	66.8	42.6	53.6	51.3	53.1	54.6
Uncertain Filter [†]	66.8	76.6	38.2	55.6	53.1	54.2	57.4
Self-Reward	66.3	76.3	45.2	55.8	56.2	54.5	59.0
WSEnsemble	64.3	74.4	44.6	54.7	55.5	54.9	58.1
AuxConf [†]	63.9	71.4	44.1	54.1	54.7	54.9	57.2
WSC Filter	69.7	<u>79.9</u>	40.5	56.4	55.8	56.1	59.7
Consultancy	61.6	68.4	44.5	54.5	54.0	52.4	55.9
Debate	55.5	51.2	46.7	48.5	51.7	49.3	50.5
SEAM w/o Filter	66.0	75.5	<u>51.0</u>	<u>56.9</u>	<u>58.5</u>	<u>56.6</u>	<u>60.7</u>
SEAM [†]	<u>67.2</u>	81.9	57.8	68.2	61.5	60.8	66.3

Table 2: Weak annotation accuracy. [†] denotes methods that require the overall distribution of the held-out data D_{held} .

Main Results

Obervation on Weak Annotation Quality Table 2 presents the annotation quality scores on different datasets. From the table, we can observe that our proposed SEAM gets a 6.5% average improvement compared to the best baseline. Besides, we have the following observations.

1) “Naïve W2S” and “Uncertain Filter” depend entirely on the weak LLMs’ capabilities for weak annotations. This indicates they struggle with annotation quality due to the limited knowledge scope of the weak LLMs, resulting in low scores for both helpfulness and harmlessness.

2) Baselines that use the confidence of pre-trained strong LLMs to correct annotations from aligned weak LLMs, such

as “AuxConf” and “WSC Filter”, exhibit extremely low weak annotation quality scores for the harmlessness objective. This occurs because the unaligned strong LLMs do not comprehend human values related to harmlessness, resulting in potential risks when their confidence is applied for corrections. In contrast, our proposed SEAM surpasses the best baseline by a significant margin on AM and CH datasets which focus on harmlessness. This highlights that our SEAM effectively improves weak annotation quality by avoiding the above risky correction issue.

Observation on W2SG Performance Table 1 presents the W2SG performance of various methods, as well as the performance of weak LLMs, and the strong ceiling (fine-tuned on ground-truth annotations). From the table, we observe that our method demonstrates an average 1.8% improvement in accuracy and a 13% increase in PGR compared to the best baseline. Besides, we can observe that:

1) W2SG is the easiest to achieve for datasets with helpfulness objectives. For instance, on the HelpSteer dataset, all baselines (except Naïve W2S) and our method achieve perfect W2SG, i.e., $PGR \geq 100\%$.

2) W2SG is very challenging for datasets with harmlessness and conflict objectives. For example, on AHarmless and SafeRLHF datasets, strong LLMs fine-tuned via all baseline W2SG methods exhibit lower performance than weak LLMs, i.e., no W2SG phenomena occurs. In contrast, our SEAM shows an accuracy improvement of 7.8% and 0.7% over the best baseline on these two datasets, respectively, and achieves W2SG with PGR values of 29% and 10%.

3) There is a positive correlation between weak annotation quality and W2SG performance. For instance, the methods with the highest and second-highest weak annotation quality scores in Table 2, i.e., SEAM and SEAM w/o filter, also display the best and second-best W2SG performance.

Method	Single				Conflict		Avg.
	Helpful		Harmless				
	AF	HS	AM	CH	AHH	SR	
SEAM	67.2	81.9	57.8	68.2	61.5	60.8	66.3
<i>w/o Filter</i>	66.0	75.5	51.0	56.9	58.5	56.6	60.7
<i>w/o Backtrace</i>	65.9	74.0	49.8	58.8	58.3	55.5	60.4
<i>w/o PAT</i>	66.4	70.2	45.6	54.5	56.7	54.3	58.0
<i>w/o FF-Tree</i>	54.5	58.4	44.8	50.6	50.3	51.4	51.7

Table 3: Quality scores of the weak annotations via various ablation versions of our method. The best results are in **bold**. Each module is removed incrementally.

Conversely, “Naïve W2S”, which has the lowest annotation quality scores, shows the worst W2SG performance. This further underscores the critical importance of improving the quality of weak annotations for enhancing W2SG.

Ablation Study

To validate the effectiveness of each strategy in our method for improving weak annotation quality, we conduct the following ablation experiments: 1) *w/o Filter*: Removing the tree information score-based filtering. 2) *w/o Backtrace*: Further removing the backtracing mechanism based on the information score in the tree search process. 3) *w/o PAT*: Further removing the principle-aware thought (PAT), meaning only the searched principle chains are provided to the weak model for annotation decisions. 4) *w/o FF-Tree*: Further removing the principle fast-and-frugal tree (FF-Tree) generation process, meaning the strong model provides sample-related knowledge without referencing principles.

From the results in Table 3. We observe that the gradual removal of each strategy progressively lowers quality scores, indicating each module’s contribution. Specifically, we find that: 1) Removing the information score-based filtering and backtracking leads to decreased annotation quality, showcasing the effectiveness of our entropy-based information score calculation in guiding filtering noise and avoiding redundancy in tree search; 2) Without knowledge derived from the strong model’s principle-aware thought, the weak model remains limited by its own knowledge scope and cannot make high-quality annotation decisions; 3) Removing the entire FF-Tree generation process leads to the most substantial decline, demonstrating its important role in eliciting useful knowledge from the strong model.

Experiments in Alignment Scenarios

To evaluate the effectiveness of our method in alignment scenarios, we treat the W2S fine-tuned M_s^{w2s} as reward models to align the pre-trained M_s . Specifically, we use prediction results (D_{eval}^{w2s}) of M_s^{w2s} on evaluation sets of AF, AM, and AHH as preference signals to align M_s using widely adopted Direct Preference Optimization (DPO) (Rafailov et al. 2024). Following Liu et al. (2024), prior to alignment, we first fine-tune M_s on the instruction-following dataset Alpaca-52K (Taori et al. 2023) to obtain




Figure 4: Evaluation results by GPT-4. We compare win/tie/lose rates of DPO-aligned strong LLMs against the SFT ones. The preference signals for DPO are from different reward models (obtained via different W2SG methods).

M_s^{sft} . Next, we perform DPO on M_s^{sft} using D_{eval}^{w2s} to obtain the aligned M_s^{dpo} . To assess the response quality of the aligned M_s^{dpo} , we select 1K prompts from SafeRLHF to generate responses using M_s^{dpo} and M_s^{sft} , which are evaluated by GPT-4 to determine the better one (following Dai et al. (2024)). Figure 4 compares the win/tie/lose rates of responses generated by different M_s^{dpo} and those generated by M_s^{sft} . From Figure 4, we can observe:

1) The alignment effect brought by the preference signals through our W2SG method is the best, far surpassing all baseline W2SG methods.

2) Surprisingly, our W2SG method has reached a strong ceiling level. Note that under the strong ceiling, the preference signals are generated by a M_s^{gt} trained on ground-truth preference annotations D_{held}^{gt} . In contrast, our method relies on the M_s^{w2s} , which is obtained through W2SG solely on unlabeled data, to generate preference signals for DPO.

3) Overall, the alignment effect obtained in the alignment scenario is positively correlated with the W2SG accuracy performance in the preference task scenario (Table 1). This further validates that exploring the W2SG problem on the preference task will directly benefit the real alignment scenario, highlighting its research value.

Conclusion

In this paper, we propose the *Strong* Empowered and *Aligned Weak* Mastered (SEAM) framework for weak annotation in Weak-to-Strong Generalization (W2SG). First, we leverage the strong model to generate a knowledge-rich principle fast-and-frugal tree to empower the annotation, alleviating the knowledge-lacking issue. Second, we position the aligned weak model as the master of the annotation process, which avoids the possible risky corrections. Empirically, we demonstrate the superiority of our method over existing W2SG baselines across six datasets, excelling in both preference tasks and alignment scenarios.

Acknowledgments

This work was supported by the grant from the National Natural Science Foundation of China (NSFC) project (No. 62276193).

References

- AI@Meta. 2024. Llama 3 Model Card. https://github.com/meta-llama/llama3/blob/main/MODEL_CARD.md.
- Bai, Y.; Jones, A.; Ndousse, K.; Askell, A.; Chen, A.; Das-Sarma, N.; Drain, D.; Fort, S.; Ganguli, D.; Henighan, T.; Joseph, N.; Kadavath, S.; Kernion, J.; Conerly, T.; El-Showk, S.; Elhage, N.; Hatfield-Dodds, Z.; Hernandez, D.; Hume, T.; Johnston, S.; Kravec, S.; Lovitt, L.; Nanda, N.; Olsson, C.; Amodei, D.; Brown, T.; Clark, J.; McCandlish, S.; Olah, C.; Mann, B.; and Kaplan, J. 2022a. Training a Helpful and Harmless Assistant with Reinforcement Learning from Human Feedback. *arXiv:2204.05862*.
- Bai, Y.; Jones, A.; Ndousse, K.; Askell, A.; Chen, A.; Das-Sarma, N.; Drain, D.; Fort, S.; Ganguli, D.; Henighan, T.; et al. 2022b. Training a helpful and harmless assistant with reinforcement learning from human feedback. *arXiv preprint arXiv:2204.05862*.
- Bai, Y.; Kadavath, S.; Kundu, S.; Askell, A.; Kernion, J.; Jones, A.; Chen, A.; Goldie, A.; Mirhoseini, A.; McKinnon, C.; et al. 2022c. Constitutional ai: Harmlessness from ai feedback. *arXiv preprint arXiv:2212.08073*.
- Bowman, S. R.; Hyun, J.; Perez, E.; Chen, E.; Pettit, C.; Heiner, S.; Lukošiušė, K.; Askell, A.; Jones, A.; Chen, A.; et al. 2022. Measuring progress on scalable oversight for large language models. *arXiv preprint arXiv:2211.03540*.
- Burns, C.; Izmailov, P.; Kirchner, J. H.; Baker, B.; Gao, L.; Aschenbrenner, L.; Chen, Y.; Ecoffet, A.; Joglekar, M.; Leike, J.; Sutskever, I.; and Wu, J. 2024. Weak-to-Strong Generalization: Eliciting Strong Capabilities With Weak Supervision. In *Forty-first International Conference on Machine Learning*.
- Cao, B.; Lu, K.; Lu, X.; Chen, J.; Ren, M.; Xiang, H.; Liu, P.; Lu, Y.; He, B.; Han, X.; et al. 2024. Towards Scalable Automated Alignment of LLMs: A Survey. *arXiv preprint arXiv:2406.01252*.
- Dai, J.; Pan, X.; Sun, R.; Ji, J.; Xu, X.; Liu, M.; Wang, Y.; and Yang, Y. 2024. Safe RLHF: Safe Reinforcement Learning from Human Feedback. In *The Twelfth International Conference on Learning Representations*.
- Dong, H.; Xiong, W.; Pang, B.; Wang, H.; Zhao, H.; Zhou, Y.; Jiang, N.; Sahoo, D.; Xiong, C.; and Zhang, T. 2024. RLhf workflow: From reward modeling to online rlhf. *arXiv preprint arXiv:2405.07863*.
- Gabriel, I. 2020. Artificial intelligence, values, and alignment. *Minds and machines*, 30(3): 411–437.
- Gigerenzer, G.; and Gaissmaier, W. 2011. Heuristic decision making. *Annual review of psychology*, 62(1): 451–482.
- Green, L.; and Mehr, D. R. 1997. What alters physicians’ decisions to admit to the coronary care unit? *Journal of Family Practice*, 45(3): 219–227.
- Guo, J.; Chen, H.; Wang, C.; Han, K.; Xu, C.; and Wang, Y. 2024a. Vision superalignment: Weak-to-strong generalization for vision foundation models. *arXiv preprint arXiv:2402.03749*.
- Guo, S.; Zhang, B.; Liu, T.; Liu, T.; Khalman, M.; Llinares, F.; Rame, A.; Mesnard, T.; Zhao, Y.; Piot, B.; et al. 2024b. Direct language model alignment from online ai feedback. *arXiv preprint arXiv:2402.04792*.
- Guo, Y.; and Yang, Y. 2024. Improving Weak-to-Strong Generalization with Reliability-Aware Alignment. *arXiv preprint arXiv:2406.19032*.
- Ildiz, M. E.; Gozeten, H. A.; Taga, E. O.; Mondelli, M.; and Oymak, S. 2024. High-dimensional Analysis of Knowledge Distillation: Weak-to-Strong Generalization and Scaling Laws. *arXiv:2410.18837*.
- Ji, J.; Qiu, T.; Chen, B.; Zhang, B.; Lou, H.; Wang, K.; Duan, Y.; He, Z.; Zhou, J.; Zhang, Z.; Zeng, F.; Ng, K. Y.; Dai, J.; Pan, X.; O’Gara, A.; Lei, Y.; Xu, H.; Tse, B.; Fu, J.; McAleer, S.; Yang, Y.; Wang, Y.; Zhu, S.-C.; Guo, Y.; and Gao, W. 2024. AI Alignment: A Comprehensive Survey. *arXiv:2310.19852*.
- Khan, A.; Hughes, J.; Valentine, D.; Ruis, L.; Sachan, K.; Radhakrishnan, A.; Grefenstette, E.; Bowman, S. R.; Rocktäschel, T.; and Perez, E. 2024. Debating with More Persuasive LLMs Leads to More Truthful Answers. In *Forty-first International Conference on Machine Learning*.
- Lee, H.; Phatale, S.; Mansoor, H.; Lu, K.; Mesnard, T.; Bishop, C.; Carbune, V.; and Rastogi, A. 2023. RLHF: Scaling reinforcement learning from human feedback with ai feedback. *arXiv preprint arXiv:2309.00267*.
- Leike, J.; Krueger, D.; Everitt, T.; Martic, M.; Maini, V.; and Legg, S. 2018. Scalable agent alignment via reward modeling: a research direction. *arXiv preprint arXiv:1811.07871*.
- Li, M.; Zhang, Y.; He, S.; Li, Z.; Zhao, H.; Wang, J.; Cheng, N.; and Zhou, T. 2024. Superfiltering: Weak-to-strong data filtering for fast instruction-tuning. *arXiv preprint arXiv:2402.00530*.
- Liu, A.; Bai, H.; Lu, Z.; Kong, X.; Wang, S.; Shan, J.; Cao, M.; and Wen, L. 2024. Direct large language model alignment through self-rewarding contrastive prompt distillation. *arXiv preprint arXiv:2402.11907*.
- Liu, Y.; and Alahi, A. 2024. Co-supervised learning: Improving weak-to-strong generalization with hierarchical mixture of experts. *arXiv preprint arXiv:2402.15505*.
- Lyu, Y.; Yan, L.; Wang, Z.; Yin, D.; Ren, P.; de Rijke, M.; and Ren, Z. 2024. MACPO: Weak-to-Strong Alignment via Multi-Agent Contrastive Preference Optimization. *arXiv:2410.07672*.
- Meng, Y.; Xia, M.; and Chen, D. 2024. Simpo: Simple preference optimization with a reference-free reward. *arXiv preprint arXiv:2405.14734*.
- Michael, J.; Mahdi, S.; Rein, D.; Petty, J.; Dirani, J.; Padmakumar, V.; and Bowman, S. R. 2023. Debate helps supervise unreliable experts. *arXiv preprint arXiv:2311.08702*.
- OpenAI. 2024a. GPT-4 Technical Report. *arXiv:2303.08774*.

- OpenAI. 2024b. Model Spec. <https://cdn.openai.com/spec/model-spec-2024-05-08.html>.
- Ouyang, L.; Wu, J.; Jiang, X.; Almeida, D.; Wainwright, C.; Mishkin, P.; Zhang, C.; Agarwal, S.; Slama, K.; Ray, A.; et al. 2022. Training language models to follow instructions with human feedback. *Advances in neural information processing systems*, 35: 27730–27744.
- Pu, X.; Gao, M.; and Wan, X. 2023. Summarization is (almost) dead. *arXiv preprint arXiv:2309.09558*.
- Rafailov, R.; Sharma, A.; Mitchell, E.; Manning, C. D.; Ermon, S.; and Finn, C. 2024. Direct preference optimization: Your language model is secretly a reward model. *Advances in Neural Information Processing Systems*, 36.
- Shin, C.; Cooper, J.; and Sala, F. 2024. Weak-to-Strong Generalization Through the Data-Centric Lens. *arXiv:2412.03881*.
- Silver, D.; Schrittwieser, J.; Simonyan, K.; Antonoglou, I.; Huang, A.; Guez, A.; Hubert, T.; Baker, L.; Lai, M.; Bolton, A.; et al. 2017. Mastering the game of go without human knowledge. *nature*, 550(7676): 354–359.
- Sun, Z.; Shen, Y.; Zhou, Q.; Zhang, H.; Chen, Z.; Cox, D.; Yang, Y.; and Gan, C. 2024. Principle-driven self-alignment of language models from scratch with minimal human supervision. *Advances in Neural Information Processing Systems*, 36.
- Tao, L.; and Li, Y. 2024. Your Weak LLM is Secretly a Strong Teacher for Alignment. *arXiv:2409.08813*.
- Taori, R.; Gulrajani, I.; Zhang, T.; Dubois, Y.; Li, X.; Guestrin, C.; Liang, P.; and Hashimoto, T. B. 2023. Stanford Alpaca: An Instruction-following LLaMA model. https://github.com/tatsu-lab/stanford_alpaca.
- Tunstall, L.; Beeching, E.; Lambert, N.; Rajani, N.; Rasul, K.; Belkada, Y.; Huang, S.; von Werra, L.; Fourrier, C.; Habib, N.; et al. 2023. Zephyr: Direct distillation of lm alignment. *arXiv preprint arXiv:2310.16944*.
- Wang, G.; Cheng, S.; Zhan, X.; Li, X.; Song, S.; and Liu, Y. 2024. OpenChat: Advancing Open-source Language Models with Mixed-Quality Data. In *The Twelfth International Conference on Learning Representations*.
- Wang, Y.; Zhong, W.; Li, L.; Mi, F.; Zeng, X.; Huang, W.; Shang, L.; Jiang, X.; and Liu, Q. 2023a. Aligning Large Language Models with Human: A Survey. *arXiv:2307.12966*.
- Wang, Z.; Dong, Y.; Zeng, J.; Adams, V.; Sreedhar, M. N.; Egert, D.; Delalleau, O.; Scowcroft, J. P.; Kant, N.; Swope, A.; and Kuchaiev, O. 2023b. HelpSteer: Multi-attribute Helpfulness Dataset for SteerLM. *arXiv:2311.09528*.
- Wu, D. X.; and Sahai, A. 2024. Provable Weak-to-Strong Generalization via Benign Overfitting. *arXiv:2410.04638*.
- Yang, A.; Yang, B.; Hui, B.; Zheng, B.; Yu, B.; Zhou, C.; Li, C.; Li, C.; Liu, D.; Huang, F.; et al. 2024a. Qwen2 technical report. *arXiv preprint arXiv:2407.10671*.
- Yang, W.; Shen, S.; Shen, G.; Gong, Z.; and Lin, Y. 2024b. Super (ficial)-alignment: Strong Models May Deceive Weak Models in Weak-to-Strong Generalization. *arXiv preprint arXiv:2406.11431*.
- Yang, Y.; Ma, Y.; and Liu, P. 2024. Weak-to-Strong Reasoning. *arXiv preprint arXiv:2407.13647*.
- Zhao, Y.; Joshi, R.; Liu, T.; Khalman, M.; Saleh, M.; and Liu, P. J. 2023. Slic-hf: Sequence likelihood calibration with human feedback. *arXiv preprint arXiv:2305.10425*.
- Zhu, W.; He, Z.; Wang, X.; Liu, P.; and Wang, R. 2024. Weak-to-Strong Preference Optimization: Stealing Reward from Weak Aligned Model. *arXiv:2410.18640*.